

# Neural code generation: course overview

---

Instructors: Sean Welleck and Daniel Fried

TAs: Nikitha Rao and Zhiruo (Zora) Wang


LTI 11891, Carnegie Mellon University, Spring 2024

[\*https://cmu-codegen.github.io/s2024\*](https://cmu-codegen.github.io/s2024)

# Sequence-to-sequence generation

## General-purpose sequence generation

- Summarize documents
- Have a conversation
- ...



## *Code* generation

- Write software
- Automatically fix bugs
- Help prove that code is correct
- Tool for reasoning
- Interact with an environment
- ...

# Code generation - applications

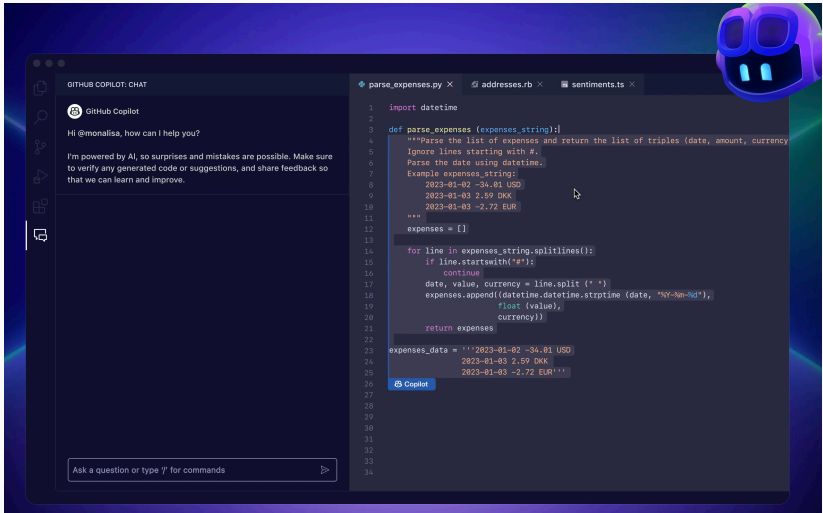


Figure 1: GitHub Copilot (12.2023)



Figure 2: FunSearch by Deepmind (12.2023)

# Code generation - applications

```
def priority(el: tuple[int, ...],  
↳ n: int) -> float:  
    score = n  
    in_el = 0  
    el_count = el.count(0)  
  
    if el_count == 0:  
        score += n ** 2  
        if el[1] == el[-1]:  
            score **= 1.5  
        if el[2] == el[-2]:  
            score **= 1.5  
        if el[3] == el[-3]:  
            score **= 1.5  
    else:  
        if el[1] == el[-1]:  
            score **= 0.5  
        if el[2] == el[-2]:  
            score **= 0.5  
  
    for e in el:  
        if e == 0:  
            if in_el == 0:  
                score **= n * 0.5  
            elif in_el == el_count - 1:  
                score **= 0.5  
            else:  
                score **= n * 0.5 ** in_el  
            in_el += 1  
        else:  
            score += 1  
  
    if el[1] == el[-1]:  
        score **= 1.5  
    if el[2] == el[-2]:  
        score **= 1.5  
  
    return score
```

**Figure 3:** The function discovered by FunSearch that results in the largest known cap set (size 512) in 8 dimensions.

Code generation with deep learning methods, primarily *neural language models*.

Example: Codex.

## Neural code generation – a brief history

Classical methods for program synthesis (specification  $\rightarrow$  program)



# Neural code generation – a brief history

Classical methods for program synthesis (specification  $\rightarrow$  program)

- *Sketch* [Solar-Lezama 2008] :

<pre>int bar(int x){     int t = x*??;     assert t == x+x;     return t; }</pre>		<pre>int bar(int x){     int t = x*2;     assert t == x+x;     return t; }</pre>
---	--	--


**Fig. 4.** Simple illustration of the integer hole.

- Specification: code with holes and test cases
- Output: fills in holes
- SAT-based search procedure

# Neural code generation – a brief history

Classical methods for program synthesis (specification → program)

- *FlashFill* [Gulwani 2011] :



- Specification: (input, output) examples
- Output: Excel string transformation
- Domain-specific language and exhaustive search

# Neural code generation – a brief history

Classical methods for program synthesis (specification  $\rightarrow$  program)

- Large search space over programs
- Difficult to model ‘informal’ specifications

# Neural code generation – a brief history

## Early language models for code

- N-gram language models [Hindle et al 2012, Allamanis & Sutton 2013]

*Programming languages, in theory, are complex, flexible and powerful, but, “natural” programs, the ones that real people actually write, are mostly simple and rather repetitive; thus they have usefully predictable statistical properties that can be captured in statistical language models and leveraged for software engineering tasks.*

Figure 4: Hindle et al 2012

# Neural code generation – a brief history

## Early language models for code

- N-gram language models [Hindle et al 2012, Allamanis & Sutton 2013]

$$p(a_4|a_1a_2a_3) = \frac{\textit{count}(a_1a_2a_3a_4)}{\textit{count}(a_1a_2a_3*)}$$

Figure 5: Hindle et al 2012

# Neural code generation – a brief history

## Early language models for code

- N-gram language models [Hindle et al 2012, Allamanis & Sutton 2013]





Figure 6: Hindle et al 2012; language-model suggestions in Eclipse

# Neural code generation – a brief history

## Early language models for code

- N-gram language models [Hindle et al 2012, Allamanis & Sutton 2013]



**Figure 6:** Hindle et al 2012; language-model suggestions in Eclipse

Restrictive n-gram model; limited generation capability

# Neural code generation – a brief history

## Early neural models for code

- Latent predictor network [Ling et al 2016]: seq2seq architecture for code generation

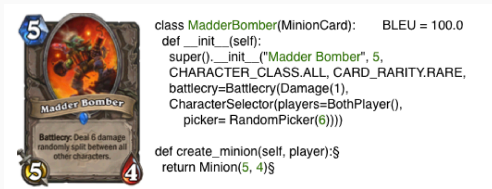


Figure 7: Generate code from a description of a card



# Neural code generation – a brief history

## Early neural models for code

- Latent predictor network [Ling et al 2016]: seq2seq architecture for code generation

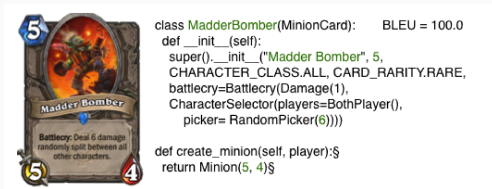


Figure 7: Generate code from a description of a card

Specialized architecture, trained for a specific dataset

## Code generation with large language models (LLMs)

---

### Evaluating Large Language Models Trained on Code

---

Mark Chen<sup>\*1</sup> Jerry Tworek<sup>\*1</sup> Heewoo Jun<sup>\*1</sup> Qiming Yuan<sup>\*1</sup> Henrique Ponde de Oliveira Pinto<sup>\*1</sup>  
Jared Kaplan<sup>\*2</sup> Harri Edwards<sup>1</sup> Yuri Burda<sup>1</sup> Nicholas Joseph<sup>2</sup> Greg Brockman<sup>1</sup> Alex Ray<sup>1</sup> Raul Puri<sup>1</sup>  
Gretchen Krueger<sup>1</sup> Michael Petrov<sup>1</sup> Heidy Khlaaf<sup>1</sup> Girish Sastry<sup>1</sup> Pamela Mishkin<sup>1</sup> Brooke Chan<sup>1</sup>  
Scott Gray<sup>1</sup> Nick Ryder<sup>1</sup> Mikhail Pavlov<sup>1</sup> Alethea Power<sup>1</sup> Lukasz Kaiser<sup>1</sup> Mohammad Bavarian<sup>1</sup>  
Clemens Winter<sup>1</sup> Philippe Tillet<sup>1</sup> Felipe Petroski Such<sup>1</sup> Dave Cummings<sup>1</sup> Matthias Plappert<sup>1</sup>  
Fotios Chantzis<sup>1</sup> Elizabeth Barnes<sup>1</sup> Ariel Herbert-Voss<sup>1</sup> William Hebgren Guss<sup>1</sup> Alex Nichol<sup>1</sup> Alex Paino<sup>1</sup>  
Nikolas Tezak<sup>1</sup> Jie Tang<sup>1</sup> Igor Babuschkin<sup>1</sup> Suchir Balaji<sup>1</sup> Shantanu Jain<sup>1</sup> William Saunders<sup>1</sup>  
Christopher Hesse<sup>1</sup> Andrew N. Carr<sup>1</sup> Jan Leike<sup>1</sup> Josh Achiam<sup>1</sup> Vedant Misra<sup>1</sup> Evan Morikawa<sup>1</sup>  
Alec Radford<sup>1</sup> Matthew Knight<sup>1</sup> Miles Brundage<sup>1</sup> Mira Murati<sup>1</sup> Katie Mayer<sup>1</sup> Peter Welinder<sup>1</sup>  
Bob McGrew<sup>1</sup> Dario Amodei<sup>2</sup> Sam McCandlish<sup>2</sup> Ilya Sutskever<sup>1</sup> Wojciech Zaremba<sup>1</sup>

#### Abstract

We introduce Codex, a GPT language model fine-tuned on publicly available code from GitHub, and study its Python code-writing capabilities.

#### 1. Introduction

Scalable sequence prediction models (Graves, 2014; Vaswani et al., 2017; Child et al., 2019) have become a general-purpose method for generation and representation learning in many domains, including natural language processing (Mikolov et al. 2013; Sutskever et al. 2014; Dai &

G] 14 Jul 2021

## Code generation with large language models (LLMs)

**Program Synthesis with Large Language Models**

Jacob Austin\*                      Augustus Odena\*

Maxwell Nye<sup>†</sup>   Maarten Bosma   Henryk Michalewski   David Dohan   Ellen Jiang   Carrie Cai

Michael Terry                      Quoc Le                      Charles Sutton

Google Research  
\* denotes equal contribution  
jaaustin@google.com, augustusodena@google.com

**Abstract**

This paper explores the limits of the current generation of large language models for program synthesis in general purpose programming languages. We evaluate a collection of such models (with between 244M and 137B parameters) on two new benchmarks, MBPP and MathQA-Python, in both the few-shot and

| 16 Aug 2021

# Neural code generation – a brief history

Code generation with large language models (LLMs)

- Language models
- + general purpose architecture
- + diverse data

# Neural code generation – a brief history

## Code generation with large language models (LLMs)

prompt

Write a function to find the smallest missing element in a sorted array. Your code should satisfy these tests:

```
assert smallest_missing([0, 1, 2, 3, 4, 5, 6], 0, 6) == 7
assert smallest_missing([0, 1, 2, 6, 9, 11, 15], 0, 6) == 3
assert smallest_missing([1, 2, 3, 4, 6, 9, 11, 15], 0, 7) == 0
```

model

```
def smallest_missing(arr, n, m):
    smallest = min(n, m)
    for i in range(n, m + 1):
        if arr[i] <= smallest:
            smallest += 1
    return smallest
```

Figure 8: Allows for natural language specifications [Austin et al 2021]

# Neural code generation – a brief history

## Code generation with large language models (LLMs)

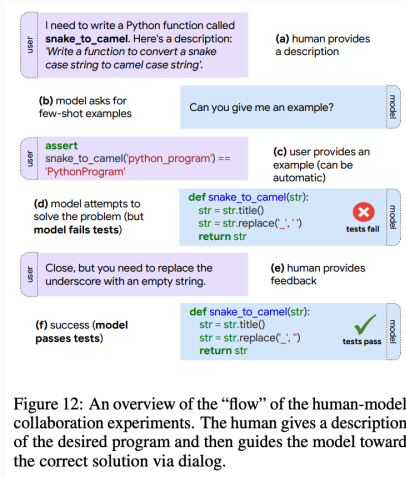


Figure 9: Key property: **flexibility** to perform many tasks [Austin et al 2021]

# Neural code generation – a brief history

## Code generation with large language models (LLMs)




Figure 10: Key property: improves by **increasing scale** [Chen et al 2021]

# Neural code generation – after Codex

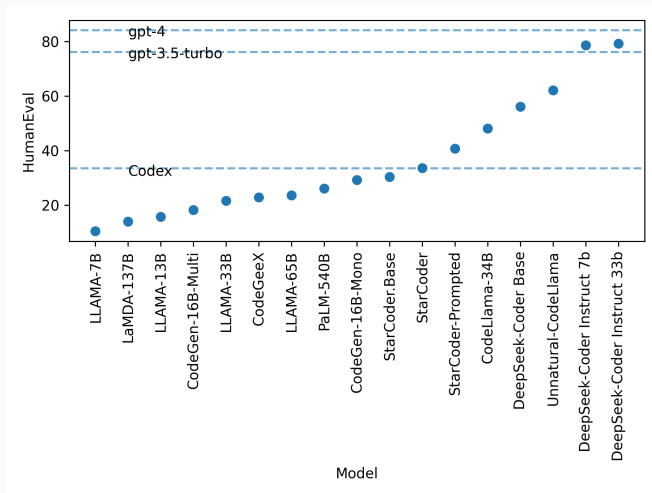


Figure 11: A lot of interest and development!



# Why neural code generation?

- Many applications
- Large amount of data
- Structured, compositional
- Combines informal (e.g., intent) and formal (e.g. testable code)
- Rich tooling (e.g., static analysis, compilers, ...)
- Often complementary to LLMs (e.g. calculator)
- ...

- Part I: *Foundations*
- Part II: *Frontiers*

Principles of neural language models as applied to code.

- Model:  $p_{\theta}(\mathbf{y}|\mathbf{x}; \mathcal{D})$ 
  - $\mathbf{x}, \mathbf{y}$  : input, output sequences
  - $\theta$  : parameters (e.g., transformer)
  - $\mathcal{D}$  : dataset

Principles of neural language models as applied to code.

- Model:  $p_{\theta}(\mathbf{y}|\mathbf{x}; \mathcal{D})$ 
  - $\mathbf{x}, \mathbf{y}$  : input, output sequences
  - $\theta$  : parameters (e.g., transformer)
  - $\mathcal{D}$  : dataset
- Learning:
  - $\arg \max_{\theta} \sum_{\mathbf{y} \in \mathcal{D}} \log p_{\theta}(\mathbf{y})$

Principles of neural language models as applied to code.

- Model:  $p_{\theta}(\mathbf{y}|\mathbf{x}; \mathcal{D})$ 
  - $\mathbf{x}, \mathbf{y}$  : input, output sequences
  - $\theta$  : parameters (e.g., transformer)
  - $\mathcal{D}$  : dataset
- Learning:
  - $\arg \max_{\theta} \sum_{\mathbf{y} \in \mathcal{D}} \log p_{\theta}(\mathbf{y})$

Principles of neural language models as applied to code.

- Model:  $p_{\theta}(\mathbf{y}|\mathbf{x}; \mathcal{D})$ 
  - $\mathbf{x}, \mathbf{y}$ : input, output sequences
  - $\theta$ : parameters (e.g., transformer)
  - $\mathcal{D}$ : dataset
- Learning:
  - $\arg \max_{\theta} \sum_{\mathbf{y} \in \mathcal{D}} \log p_{\theta}(\mathbf{y})$
- Inference:
  - $\mathbf{y} = f(p_{\theta}(\cdot|\mathbf{x}))$
  - $f$ : e.g., sampling

Principles of neural language models as applied to code.

- Model:  $p_{\theta}(\mathbf{y}|\mathbf{x}; \mathcal{D})$ 
  - $\mathbf{x}, \mathbf{y}$ : input, output sequences
  - $\theta$ : parameters (e.g., transformer)
  - $\mathcal{D}$ : dataset
- Learning:
  - $\arg \max_{\theta} \sum_{\mathbf{y} \in \mathcal{D}} \log p_{\theta}(\mathbf{y})$
- Inference:
  - $\mathbf{y} = f(p_{\theta}(\cdot|\mathbf{x}))$
  - $f$ : e.g., sampling
- Evaluation

*Learning*: how do we *train* language models for code generation?

- **Pretraining**: large-scale initial training based on *scaling laws* (1/18) and *code objectives* (1/23)



*Learning*: how do we *train* language models for code generation?

- **Pretraining**: large-scale initial training based on *scaling laws* (1/18) and *code objectives* (1/23)
- **Finetuning**: specializing the model for specific tasks and languages (1/25)

*Learning*: how do we *train* language models for code generation?

- **Pretraining**: large-scale initial training based on *scaling laws* (1/18) and *code objectives* (1/23)
- **Finetuning**: specializing the model for specific tasks and languages (1/25)
- **Learning from feedback**: improving the model with feedback on its outputs, such as execution results and language (1/30)

*Evaluation*: how good is our neural code generator?

- Code metrics and benchmarks (2/01, 2/06)

*Data*: what data should we train with? (2/08, 2/13)

- Data for *pretraining* and *domain-adaptation*
- *Synthetic* data
- Impact of data *quality*

*Inference*: how do we *generate* code with a trained language model?

- Algorithms that leverage *execution*, *verification*, and *feedback*  
(2/15, 2/20)

- Part I: *Foundations*
  - Learning, Inference, Data, Evaluation
- Part II: *Frontiers*

- Part I: *Foundations*
- Part II: *Frontiers*

Code is **communicative** and code generators are **used by real people**

- **Pragmatic** aspects of code generation (2/29)
- **Programming with AI** (3/12) and dealing with **uncertainty** (3/14)
- **Guest lecture** by Sherry Wu (3/21)



Real-world code is long, exists in repositories unseen during training, and evolves over time. How do we adapt to these conditions?

- Methods for **long-context** generation and **retrieval** in code (3/19, 3/26)

## Part II: Frontiers – Reasoning

### Code as a medium for reasoning and control (4/02)

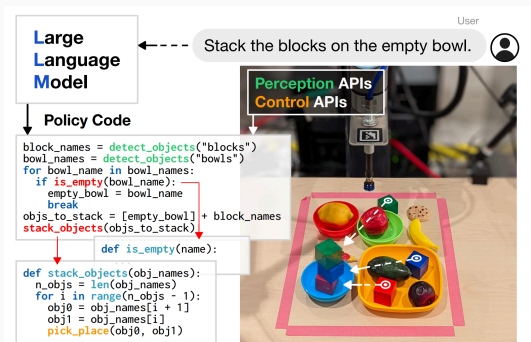


Fig. 1: Given examples (via few-shot prompting), robots can use code-writing large language models (LLMs) to translate natural language commands into robot policy code which process **perception** outputs, parameterize **control** primitives, recursively generate code for **undefined** functions, and generalize to new tasks.

Figure 12: Code generation for robotics

Some programming languages allow for **proving** that code is **correct**<sup>1</sup>

- **Neural theorem proving** (4/04)
  - Use LLMs to make it easier to verify things
  - Use verifiable code for mathematical reasoning
- Formally **verified code synthesis** (4/09)
- Guest lecture by Zhangir Azerbayev (4/18)

---

<sup>1</sup>E.g., Coq, Dafny, F\*, Isabelle, Lean

- Programs are structured, testable, interpretable.
- These properties can be leveraged by *large-scale neural program search* to **discover solutions** to open problems (4/16)

- Part I: *Foundations*
  - Learning, inference, data, evaluation
- Part II: *Frontiers*
  - Interaction, adaptability, reasoning, formal methods, science

## Course structure, projects, and logistics

---

- 6-unit version of the course
  - Attend lectures (with pre- and post-assignments)
  - Attend discussions (with pre- and post-assignments)
  - Lead a discussion with a team (via a presentation)

# Course structure

- 6-unit version of the course
  - Attend lectures (with pre- and post-assignments)
  - Attend discussions (with pre- and post-assignments)
  - Lead a discussion with a team (via a presentation)
- 12-unit version of the course: all the above, plus:
  - A high-quality research project, in teams of 2–4.
    - Two checkpoint reports
    - Two structured project hours
    - Final presentation
    - Final report



## 6-Unit course structure: discussions

In a student-led discussion, 3 students present a (set of) papers on a theme. Choose how much to focus on each paper, but cover the following topics:

- **Content:** motivation, setting, methods, findings. What was surprising?
- **Reviewer:** role-play a conference reviewer. Score the paper, and justify.
- **Future:** Brainstorm future work ideas for discussion.
- **Reproducibility:** What code and data would you use to dig deeper?

Use slides, but a main goal is to facilitate a discussion!

## 6-Unit course structure: discussions

For presenters:

- Submit your slides before the day you present.
- We'll grade based on the presentation and slides.
- It's ok if you spark a long discussion and don't get through all slides.
- Present one time during the course, for 33% of the 6-unit grade, or 16% of the 12-unit grade.

## 6-Unit course structure: discussions

For presenters:

- Submit your slides before the day you present.
- We'll grade based on the presentation and slides.
- It's ok if you spark a long discussion and don't get through all slides.
- Present one time during the course, for 33% of the 6-unit grade, or 16% of the 12-unit grade.

Sign-ups:

- Sign-up link coming after class.
- **Please sign up by Thursday end-of-day.** You can swap later if you find someone willing to.
- Extra credit (+2 out of 20 presentation points) for any team that presents on Thursday next week (1/25), on *finetuning for code*.

On days you're not presenting (**both lectures and discussions**):

- Pre-assignment (33% of grade):
  - Short summary and  $\geq 1$  discussion questions for a paper.
  - Submit by 11:59pm the day **before** class.
  - 23 days, but we'll grade out of 20.

On days you're not presenting (**both lectures and discussions**):

- Pre-assignment (33% of grade):
  - Short summary and  $\geq 1$  discussion questions for a paper.
  - Submit by 11:59pm the day **before** class.
  - 23 days, but we'll grade out of 20.
- Post-assignment (33% of grade):
  - 2-3 sentences on what you found interesting.
  - Submit by 11:59pm the day **of** class.
  - 23 days, but we'll grade out of 20.

## 12-Unit: Course project

- For students taking the class for 12 units, all of the 6 unit requirements, and also a course project.
- **Simulates doing a research project** on a topic related to the course.
- Teams of 2-4 members
- *Propose your own topic or pick a topic from our list*
- Ends in a report and presentation that should be in the style of a workshop paper or the first draft of a conference paper.

## 12-Unit: Project timeline (tentative)

- Team formation: Jan 30th

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).



## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.
- **Report 1 (25%):** Feb 30th  
Task proposal and data analysis; related work; baseline proposal.

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.
- **Report 1 (25%):** Feb 30th  
Task proposal and data analysis; related work; baseline proposal.
- **Spring break:** Mar 5th and 7th – Take time off!

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.
- **Report 1 (25%):** Feb 30th  
Task proposal and data analysis; related work; baseline proposal.
- **Spring break:** Mar 5th and 7th – Take time off!
- **Project hours 2 (5%):** Mar 28th

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.
- **Report 1 (25%):** Feb 30th  
Task proposal and data analysis; related work; baseline proposal.
- **Spring break:** Mar 5th and 7th – Take time off!
- **Project hours 2 (5%):** Mar 28th
- **Report 2 (25%):** Mar 30th  
Baseline results and analysis, and a technique proposal.

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.
- **Report 1 (25%):** Feb 30th  
Task proposal and data analysis; related work; baseline proposal.
- **Spring break:** Mar 5th and 7th – Take time off!
- **Project hours 2 (5%):** Mar 28th
- **Report 2 (25%):** Mar 30th  
Baseline results and analysis, and a technique proposal.
- **Final presentations (10%):** Apr 23rd and 25th  
In-class 15-20 minute presentations.

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

## 12-Unit: Project timeline (tentative)

- **Team formation:** Jan 30th
- **Project hours 1 (5%):** Feb 22nd  
Meet with an instructor for 15-20 minutes, with a few slides.
- **Report 1 (25%):** Feb 30th  
Task proposal and data analysis; related work; baseline proposal.
- **Spring break:** Mar 5th and 7th – Take time off!
- **Project hours 2 (5%):** Mar 28th
- **Report 2 (25%):** Mar 30th  
Baseline results and analysis, and a technique proposal.
- **Final presentations (10%):** Apr 23rd and 25th  
In-class 15-20 minute presentations.
- **Final report (30%):** Apr 29th  
Results and analysis of your technique; future work proposal.

Your team will have a total of 5 late days which you can budget across any of the written reports (Report 1, Report 2, or the Final Report).

- Introduce yourself! Name and program.
- What brings you to this class?

- Part I: *Foundations*
  - *Learning*, inference, data, evaluation
- Part II: *Frontiers*
  - Interaction, adaptability, reasoning, formal methods, science

*Next meeting*: lecture on pretraining and scaling laws for code



